

HW 2

This assignment covers several aspects of Linear Regression. **DO NOT ERASE MARKDOWN CELLS AND INSTRUCTIONS IN YOUR HW submission**

- **Q** - QUESTION
- **A** - Where to input your answer

Instructions

Keep the following in mind for all notebooks you develop:

- Structure your notebook.
- Use headings with meaningful levels in Markdown cells, and explain the questions each piece of code is to answer or the reason it is there.
- Make sure your notebook can always be rerun from top to bottom.
- Please start working on this assignment as soon as possible. If you are a beginner in Python this might take a long time. One of the objectives of this assignment is to help you learn python and scikit-learn package.
- Follow [README.md](#) for homework submission instructions

Tutorials

- [scikit-learn linear model](#)
- [train-test-split](#)
- [least squares fitting](#)
- [Linear Regression](#)
- [Seaborn](#)

REGRESSION TASK USING SKLEARN

In jupyter notebook environment, commands starting with the symbol % are magic commands or magic functions. `%%timeit` is one of such function. It basically gives you the speed of execution of certain statement or blocks of codes.

```
In [167... import pandas as pd
import numpy as np
import seaborn as sns
```

Data Get the exploratory data and the following files:

<https://archive.ics.uci.edu/ml/machine-learning-databases/auto-mpg/auto-mpg.data>

<https://archive.ics.uci.edu/ml/machine-learning-databases/auto-mpg/auto-mpg.names>

or Use from our 2024Spring/data repository folder

- Link should automatically download the data
- copy them in your HW folder
- If you are using command line: `>> wget`
`https://archive.ics.uci.edu/ml/machine-learning-databases/auto-mpg/auto-mpg.data`
 - If wget is not working
 - download it from [link](#)
 - follow [steps](#)

Q1 Read the data using pandas, and replace the ??? in the code cell below to accomplish this task. Note that auto-mpg.data does not have the column headers. use auto-mpg.names file to provide column names to the dataframe.

A1

```
In [168... # Replace ??? with code in the code cell below
column_names = ['mpg', 'cylinders', 'displacement', 'horsepower', 'weight',
df = pd.read_csv('auto-mpg.data', delim_whitespace = True, names=column_name
```

```
In [169... # View head of the data to confirm the correctness of your answer
df.head()
```

Out[169]:

	mpg	cylinders	displacement	horsepower	weight	acceleration	model_year	origin	ca
0	18.0	8	307.0	130.0	3504.0	12.0	70	1	c
1	15.0	8	350.0	165.0	3693.0	11.5	70	1	
2	18.0	8	318.0	150.0	3436.0	11.0	70	1	p
3	16.0	8	304.0	150.0	3433.0	12.0	70	1	al
4	17.0	8	302.0	140.0	3449.0	10.5	70	1	

Data cleaning and manipulation

Use

Q2 Data cleaning and manipulation:

1. use `pandas.info()` method to find columns with large number of NaN values
2. remove the column with NaN values
3. Check if there are still NaN values in the dataframe using `isna()` method

A2 Replace ??? with code in the code cell below

In [170...

```

#1. use pandas.info() method to find columns with large number of NaN values
df.info()

#2. remove the column with NaN values - replace ??? with code
df.drop(columns=['horsepower'])
# Print head
df.head()

#3. Check if there are still NaN values in the dataframe using ``isna()``
df.isna().sum()
# drop if any left or replace Nan values
df.dropna(inplace=True)

```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 398 entries, 0 to 397
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   mpg              398 non-null    float64
1   cylinders        398 non-null    int64
2   displacement     398 non-null    float64
3   horsepower       392 non-null    float64
4   weight           398 non-null    float64
5   acceleration     398 non-null    float64
6   model_year       398 non-null    int64
7   origin           398 non-null    int64
8   car_name         398 non-null    object
dtypes: float64(5), int64(3), object(1)
memory usage: 28.1+ KB
```

```
In [171]: #Print Tail
df.tail()
```

```
Out[171]:
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	model_year	origin
393	27.0	4	140.0	86.0	2790.0	15.6	82	1
394	44.0	4	97.0	52.0	2130.0	24.6	82	2
395	32.0	4	135.0	84.0	2295.0	11.6	82	1
396	28.0	4	120.0	79.0	2625.0	18.6	82	1
397	31.0	4	119.0	82.0	2720.0	19.4	82	1

Q3:

1. Convert following columns 'cylinders', 'year', 'origin' to dummy variable using pandas get_dummies() function
2. Do data normalization on real value/continuous columns
 - The formula for normalization is: $(\text{Col_value} - \text{Mean of the col}) / \text{Standard Deviation of the col}$

A3 Replace ??? with code in the code cell below

```
In [172... # 1. Convert following columns 'cylinders', 'year', 'origin' to dummy variables
cols = ['cylinders', 'model_year', 'origin']
df_dummies = pd.get_dummies(df, columns=cols)

#show the head
df_dummies.head()

# 2. Do data normalization on real value/continuous columns
realcols = ['mpg', 'displacement', 'horsepower', 'weight', 'acceleration']

for col in realcols:
    mean = df[col].mean()
    std = df[col].std()
    df[col] = (df[col] - mean) / std
```

Regression Task


Given all the information we will try to predict mpg - miles per gallon. The First step toward predicting the mpg from the dataset is to find the correlation between the columns/features of the dataset.

Q4

1. Use heatmap chart from seaborn library to findout the correlation between the columns.
2. Which of the columns is mostly related to mpg column and why?

```
In [173... # A4 code goes below
sns.heatmap(df.select_dtypes(include=[np.number]).corr(), annot=True, cmap='
```

Out[173]: <Axes: >



A4

The column that is most closely related to the mpg column is weight. This is because the weight column has the highest negative correlation with mpg at -0.83, indicating that as a car's weight increases, its mpg tends to decrease.


Q5

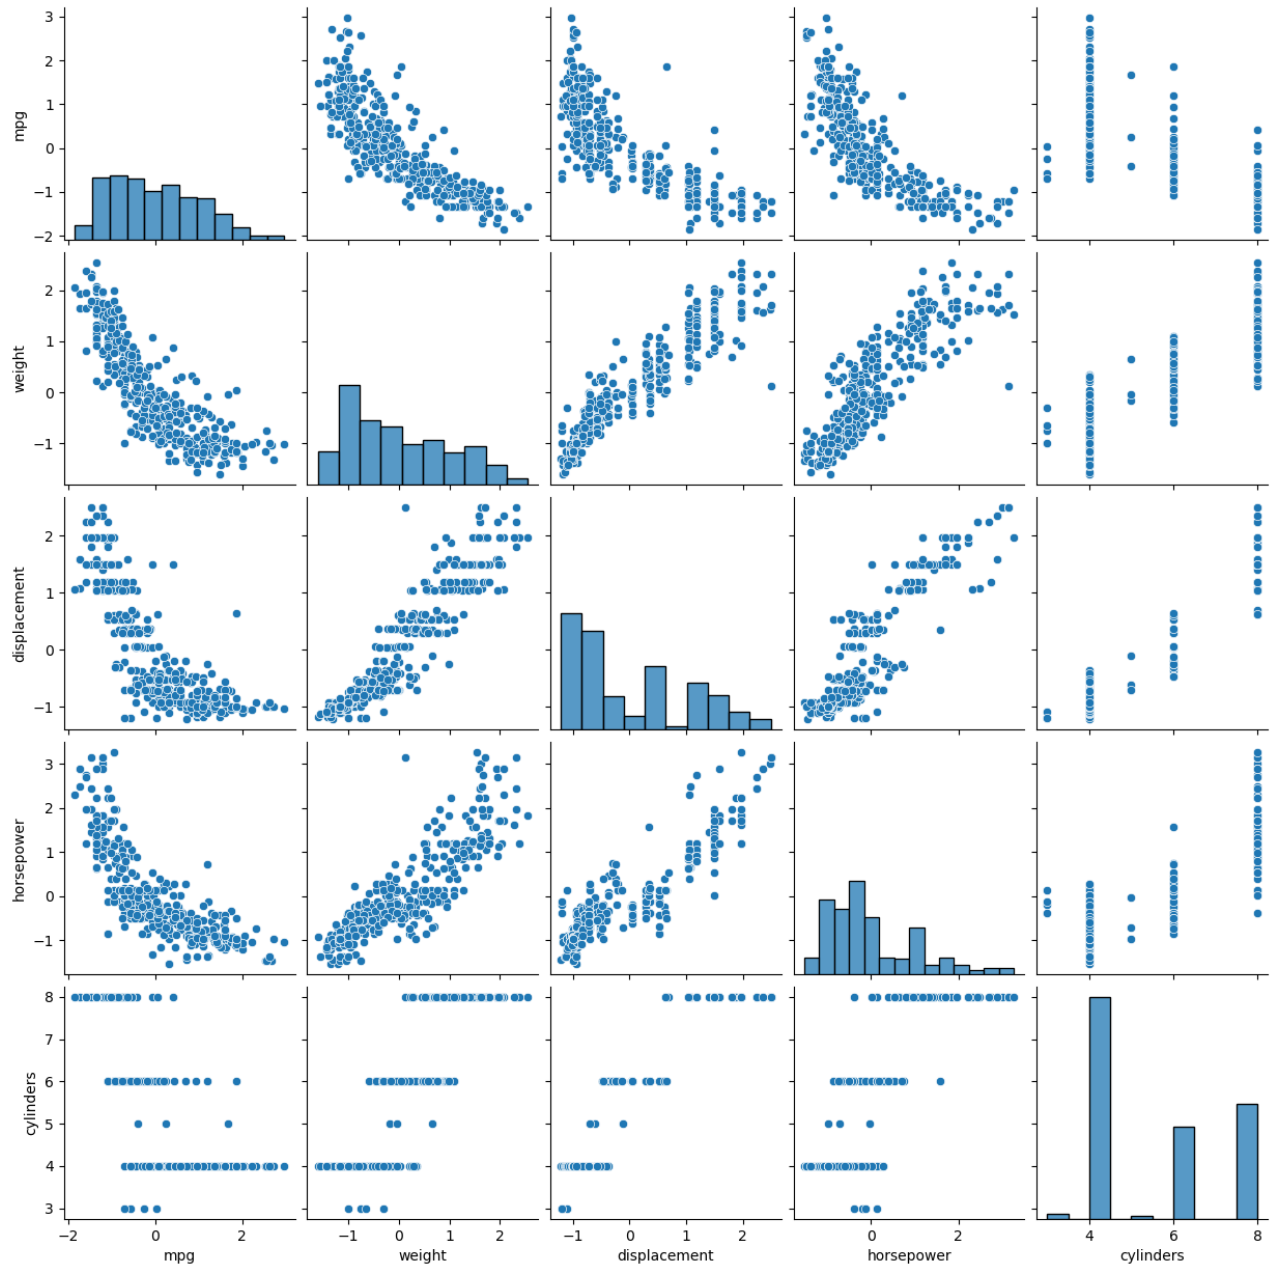
1. Draw a lineplot or scattered plot between mpg and your answer from the above cell.
 2. Use pairplot from sns to plot our data frame df for better understanding of your selection
 - NOTE: 2. should inform 3.
 3. Choose a set of columns/ features based on pairplot and heatmap for the mpg prediction.
- Justify your answer using some explanation from the heatmap and pairplot graph formulated from the dataset.

A5 For 1. and 2. replace ??? with code in the code cell below.

```
In [174... # 1. Draw a lineplot or scattered plot between mpg and your answer from the
sns.scatterplot(x='weight', y='mpg', data=df)
# 2. Use pairplot from sns to plot our data frame df for better understanding
sns.pairplot(df[['mpg', 'weight', 'displacement', 'horsepower', 'cylinders']])

/Users/javier/anaconda3/lib/python3.11/site-packages/seaborn/axisgrid.py:11
8: UserWarning: The figure layout has changed to tight
  self._figure.tight_layout(*args, **kwargs)
Out[174]: <seaborn.axisgrid.PairGrid at 0x115869450>
```





A5


Based on the visual and statistical analysis, the chosen features for predicting mpg are those that show the strongest correlations, both in the heatmap and in the pairplot. These include weight, displacement, horsepower, and cylinders. Each of these is known to have a direct impact on the vehicle's fuel efficiency, which justifies their selection for the prediction model.

Q6 Data Visualization:

1. Now, create a histogram which represents number items with per cylinder class

A6 Replace ??? with code in the code cell below

```
In [175... import matplotlib.pyplot as plt
plt.hist(df['cylinders'])
plt.show()
```



Data Preparation

Q7 Assign mpg column value to y and rest columns to x, remember x shouldn't have mpg

A7 Replace ??? with code in the code cell below

```
In [176... df = df.drop('car_name', axis=1)
y = df['mpg']
df.drop(['mpg'], axis = 1)
x = df
```

Q8 Use `train_test_split` to split the data set as train:test=(80%:20%) ratio.

A8 Replace ??? with code in the code cell below

```
In [177... from sklearn.model_selection import train_test_split

xtrain, xtest, ytrain, ytest = train_test_split(x, y, test_size=0.20)
# View the shape of your data set
xtrain.shape, xtest.shape, ytrain.shape, ytest.shape
```

```
Out[177]: ((313, 8), (79, 8), (313,), (79,))
```

Q9 Follow examples from references given in the top of this notebook

- Note: Use linear model to fit regression line and plot
- Our linear model will be of following type
- $Y = b + \text{coef}_0x_0 + \text{coef}_1x_1 + \text{coef}_2x_2 + \dots$

A9: Replace ??? with code in the code cell below

```
In [178... from sklearn import linear_model

reg = linear_model.LinearRegression()
reg.fit(xtrain, ytrain)

#Now view the coefficient use .coef_ and shape of .coef_
print(reg.coef_)
print(reg.intercept_)
print(reg.coef_.shape[0])

[ 1.00000000e+00 -1.19169384e-16 -2.14162869e-16 -3.13131711e-17
 -3.21094058e-17 -2.29479929e-16 -3.98455463e-17 -1.50130708e-16]
3.937822290467352e-15
8
```

Q10 Relates to the code in the cell below. Why the printed values the same?

```
In [179... # Now if you view
print(f'{reg.coef_.shape[0]},{xtrain.shape[1]}, ', f'are equal? {reg.coef_.s

8,8, are equal? True
```

A10 It confirms that the model structures are consistent with the data it was trained on.

Model Scoring

```
In [180... # Model Score
from sklearn import linear_model
reg = linear_model.LinearRegression()
reg.fit(xtrain, ytrain)
reg.score(xtest, ytest)

# Calculate the score on train and test sets
# Your code goes below
reg.score(xtrain, ytrain), reg.score(xtest, ytest)
```

Out[180]: (1.0, 1.0)

Q11 Each of the sklearn models have different model evaluations core value.

- LinearRegression [documentation](#)
- More on [model_evaluation](#)

Explain what's the meaning of reg.score return value in this notebook.

A11 It means that the regression predictions perfectly fit the data.

```
In [181... # A custom function to calculate r2 score
# Details on the custom scorers: https://scikit-learn.org/stable/modules/mod

def r2score_(ytrue, ypred):
    rss = ((ytrue - ypred)**2).sum()
    tss = ((ytrue - ytrue.mean()) ** 2).sum()
    r2 = 1 - rss/tss
    return r2

# Now do prediction on xtrain and xtest and check your r2 score by printing
trainpredict = reg.predict(xtrain)
testpredict = reg.predict(xtest)

print(r2score_(ytrain, trainpredict), r2score_(ytest, testpredict))
```

1.0 1.0

One way of achieving linear regression is by minimizing the error between actual y and predicted y. The method is known as least square method. We will make our custom least square optimize to calculate model parameters that minimizes output error.

Q12 Write a function which takes weights(or params), x and y and do following

- 1. calculate dot product between x and params , which is ypredicted
- 1. calculate difference between actual y and ypredicted
- 1. return the difference

A12 complete the code below

In [182...

```
import scipy.optimize as optimization
from sklearn.metrics import r2_score

def constraint(params, x, y):
    ypred = x@params
    return y-ypred

# Our initial params is a vector of size equal to dimension of x, or you can
# You can create zeros vector using np.zeros(size)

# complete code
params = np.zeros(xtrain.shape[1])

# Now study the documentation and complete following code
params, _ = optimization.leastsq(constraint, params, args=(xtrain, ytrain))

# Now we have parameter or weight we can now create our model
model = lambda x:np.dot(x,params)

# Now predict ytrain using model and see first 5 predicted and actual values
ypred_train = model(xtrain)
# see first 5 predicted values
print(ypred_train[:5])
# see first 5 actual values
print(ytrain[:5])

# Now predict ytest using model and see first 5 predicted and actual values
ypred_test = model(xtest)
print(ypred_test[:5])
print(ytest[:5])

# Now use custom made r2score calculator to calculate r2 score on both train
print(r2score_(ytest, ypred_test), r2score_(ytrain, ypred_train))

# Now use sklearn build-in r2score calculator to calculate r2 score on both
print(r2_score(ytrain, ypred_train), r2_score(ytest, ypred_test))
```

```
[ 1.39065615  1.07034896 -1.33836109  1.1344104   0.83972778]
317      1.390656
303      1.070349
105     -1.338361
346      1.134410
239      0.839728
Name: mpg, dtype: float64
[-1.21023822 -1.08211534 -0.69774672  0.58348203  0.48098373]
139     -1.210238
115     -1.082115
0        -0.697747
146      0.583482
338      0.480984
Name: mpg, dtype: float64
1.0 1.0
1.0 1.0
```